

# Introduction to Deep Learning wit R

Gabriela de Queiroz

Data Scientist and Founder of R-Ladies



[k-roz.com](http://k-roz.com)




@gdequeiroz

# What is deep learning?

- Subfield of machine learning
- New take on learning representations from data: puts an emphasis on learning **successive layers** of increasingly meaningful representations
- "deep" stands for the idea of successive layers of representations
- **Depth of the model:** how many layers contribute to a model of the data

# What is deep learning?


- In deep learning, these layered representations are (almost always) learned via models called **neural networks**, structured in literal layers stacked on top of each other.



# What is deep learning?

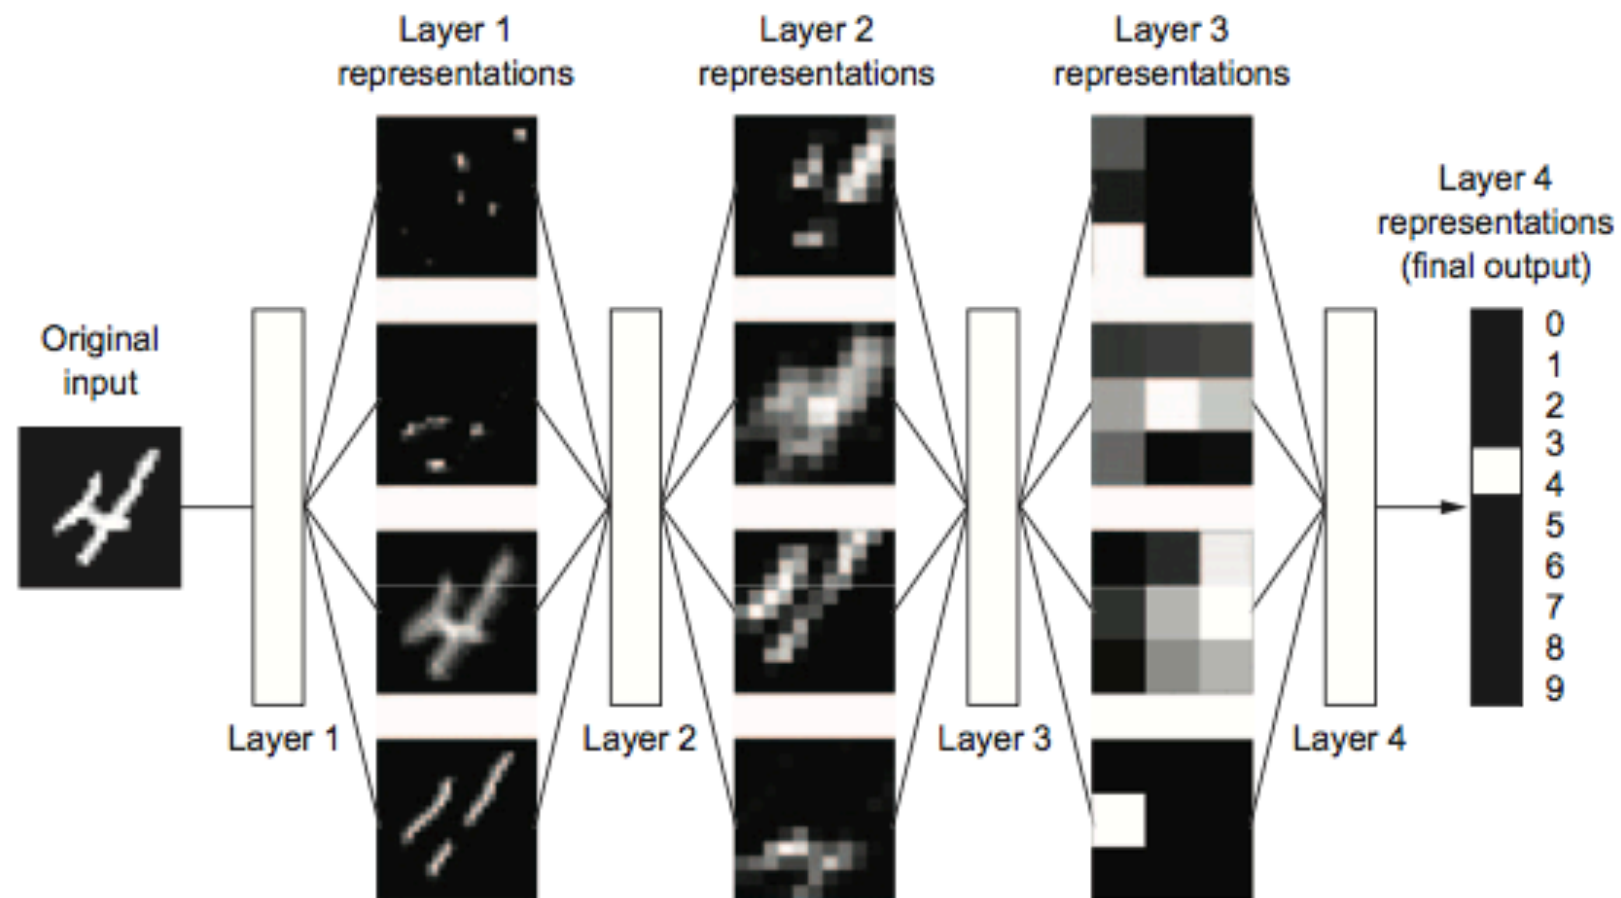
- Take some input (observations,  $X$ ) and transform into some output (predictions,  $Y$ ) via successive layers of representation

**Input**  
(grayscale  
image of a  
handwritten  
number 4)



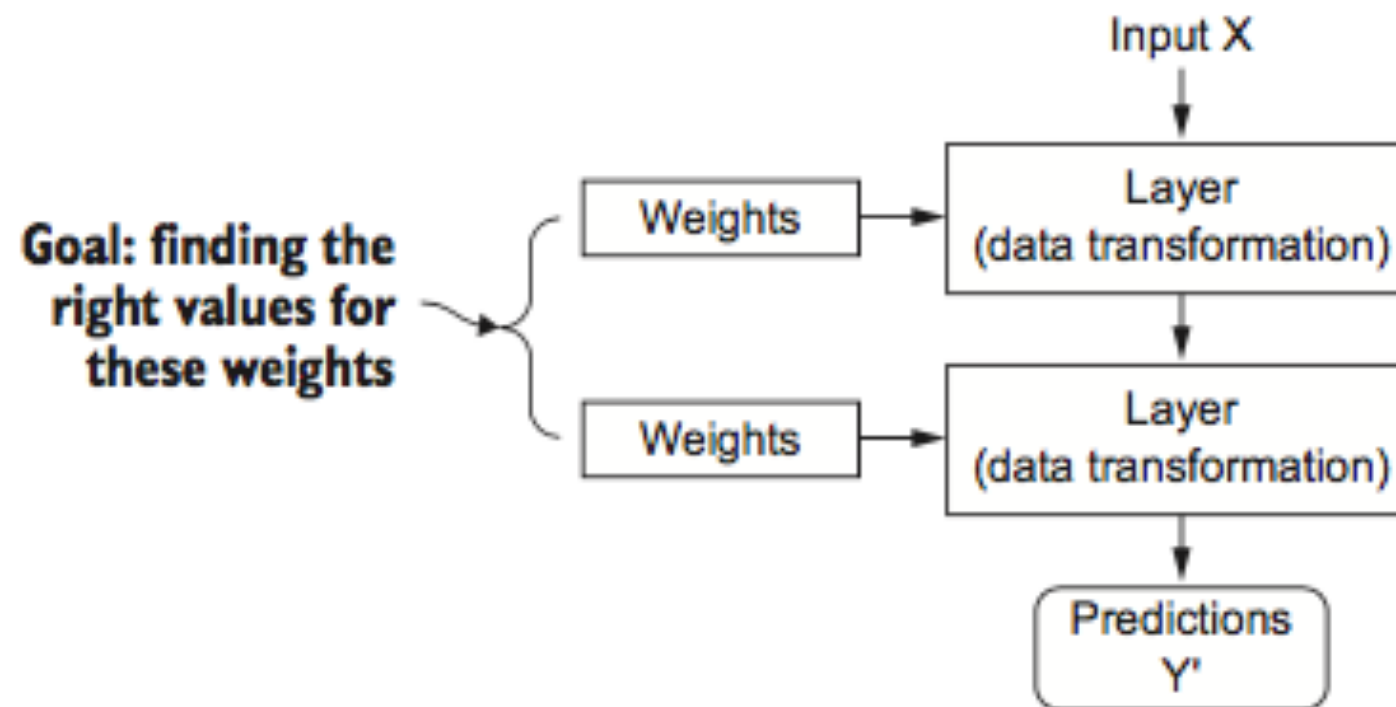
**Output**  
(prediction about  
what the digit is )

# Deep representations learned by a digit-classification model



# What are layers?

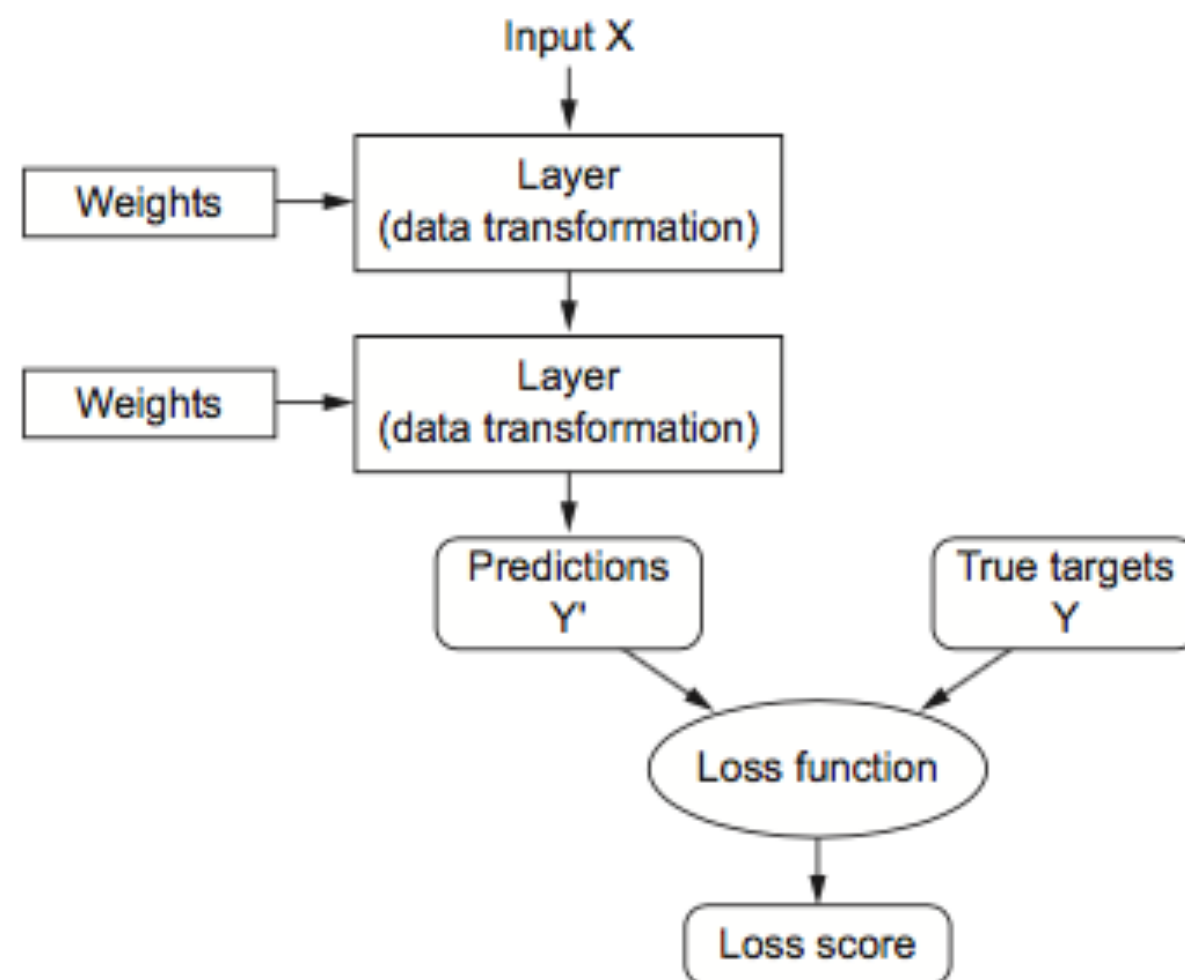
- Data transformation functions *parameterized* by weights\* (like a linear equation)



\*Weights are also sometimes called the parameters of a layer

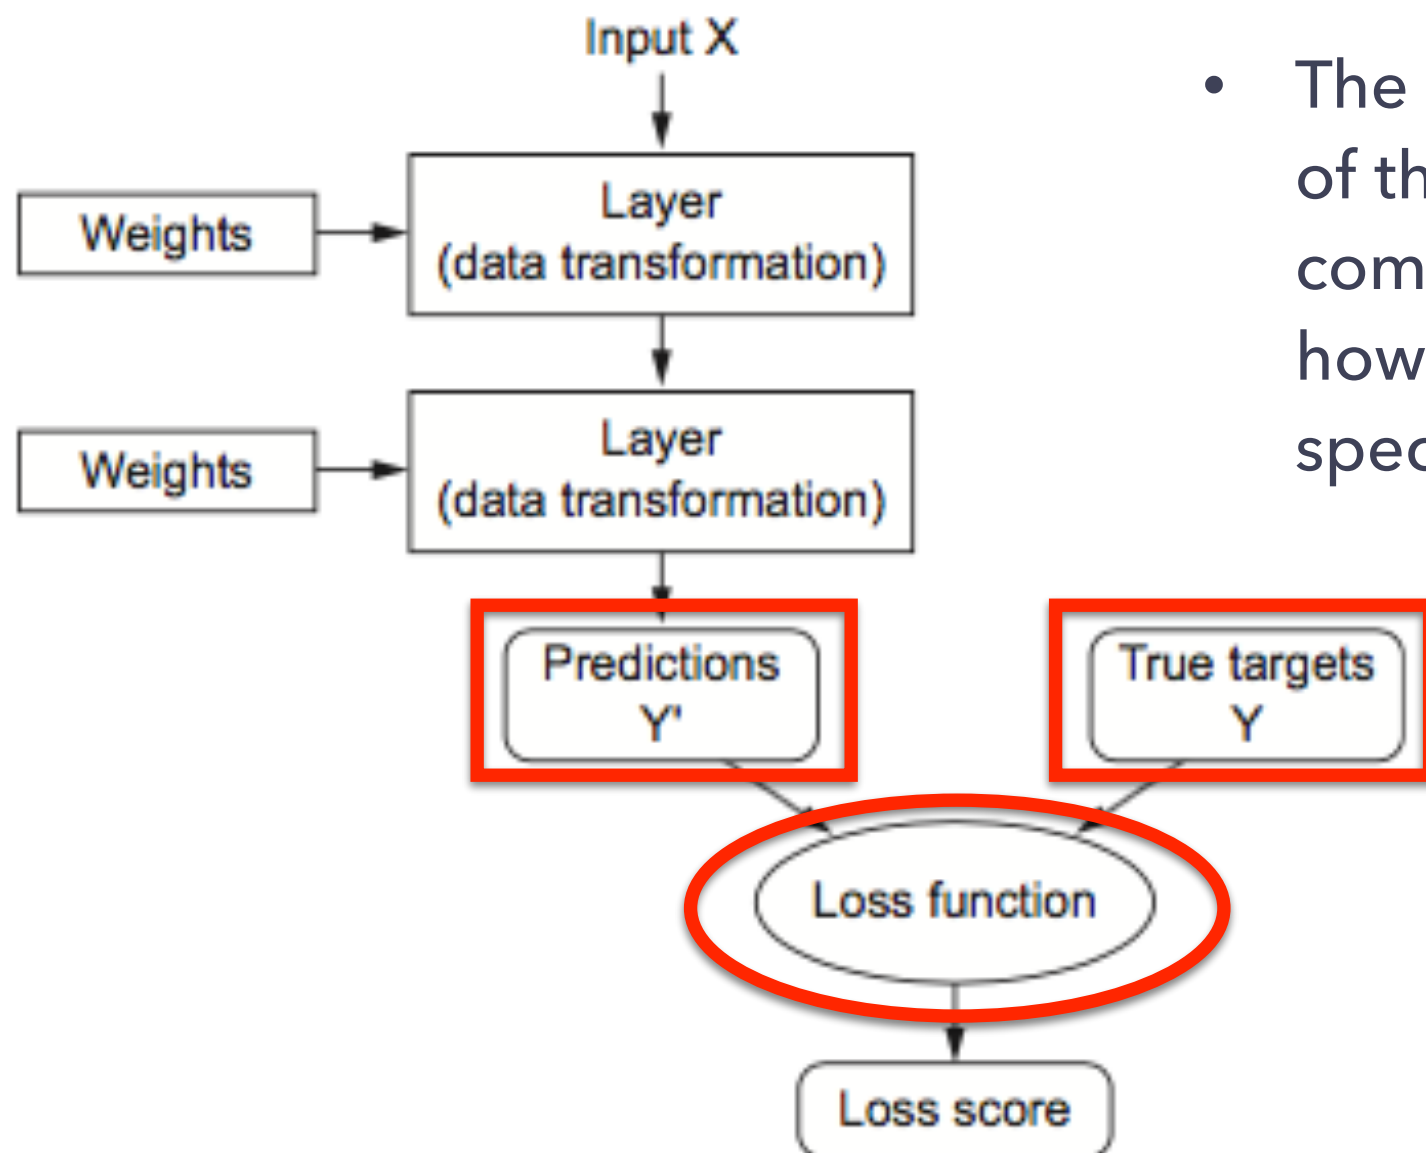
# Loss Function/ Objective Function

- To control the output of a neural network, you need to be able to measure how far this output is from what you expected (**loss function** of the network or **objective function**).



\*Weights are also sometimes called the parameters of a layer

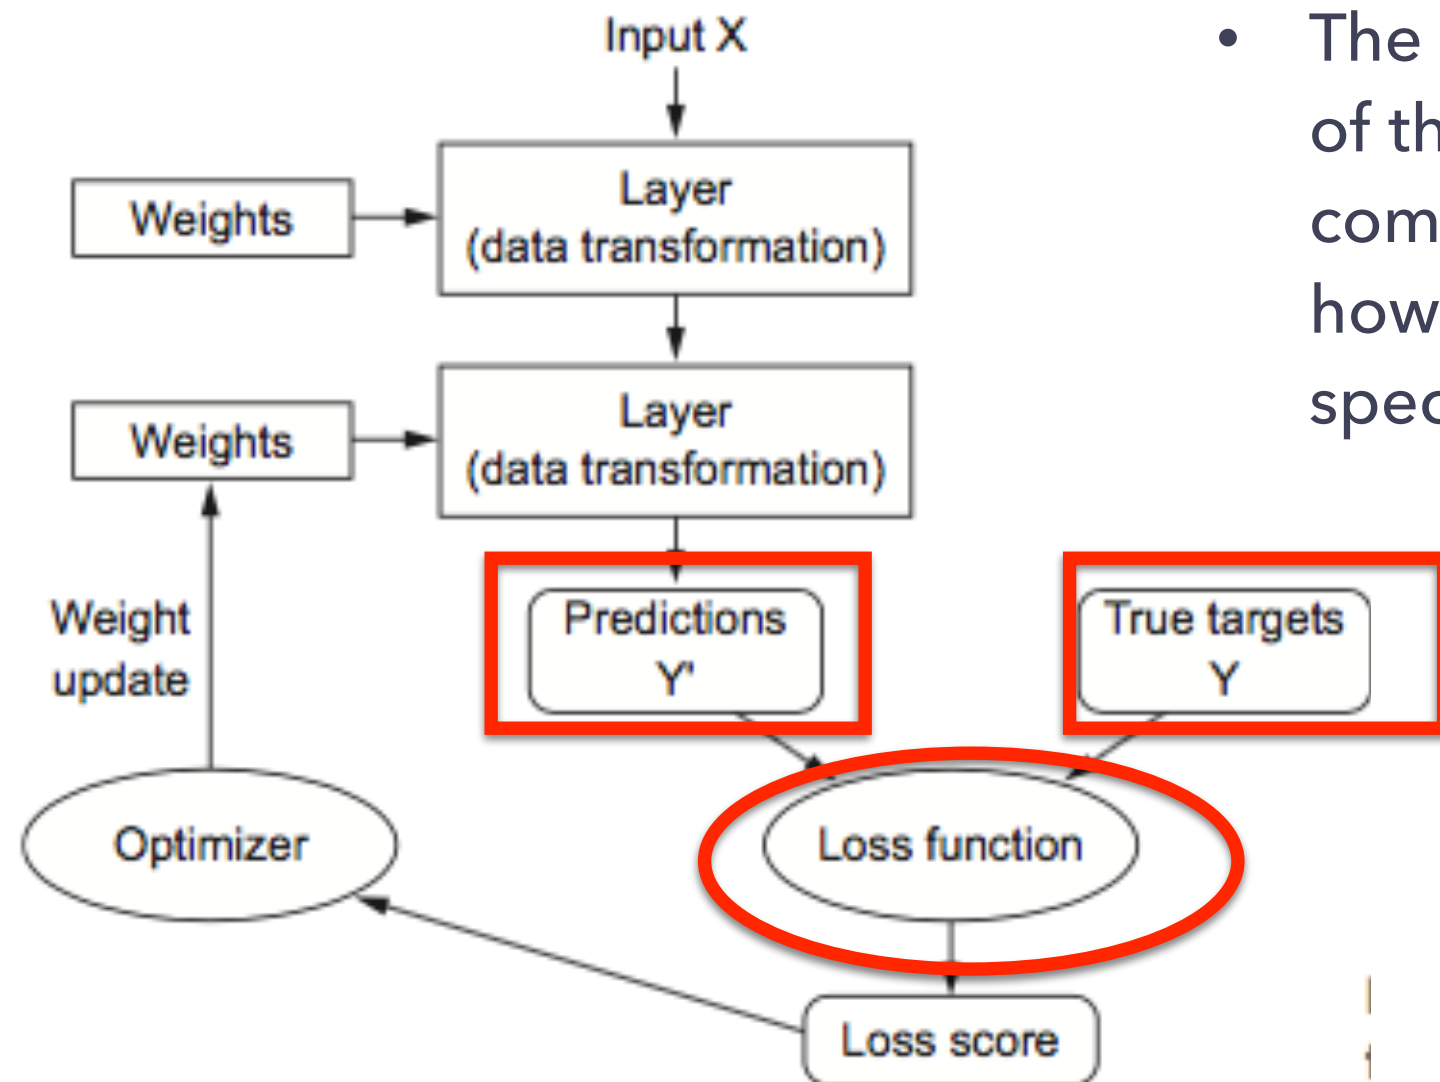
# Loss Function/ Objective Function



- The loss function takes the **predictions** of the network and the **true target** and computes a distance score, capturing how well the network has done on this specific example.



# Backpropagation Algorithm



- The loss function takes the **predictions** of the network and the **true target** and computes a distance score, capturing how well the network has done on this specific example.

**How do we train deep  
learning models?**

```
network <- keras_model_sequential() %>%  
  layer_dense(units = 512,  
              activation = "relu",  
              input_shape = c(28 * 28)) %>%  
  layer_dense(units = 10,  
              activation = "softmax")
```



10-way softmax layer, which means it will return an array of 10 probability scores (summing to 1)

# Data Representation for Neural Networks

- Tensors: generalization of vectors and matrices to an arbitrary number of dimensions (note that in the context of tensors, a dimension is often called an axis).
- In R, **vectors** are used to create and manipulate 1D tensors, and **matrices** are used for 2D tensors. For higher-level dimensions, array objects (which support any number of dimensions) are used.
-

# Three key attributes of a tensor

- **Number of axes (rank)**
- **Shape**
- **Data Types**

# Number of axes (rank)

- For instance, a 3D tensor has three axes, and a matrix has two axes.

# Shape

- **Shape:** integer vector that describes how many dimensions the tensor has along each axis.

- For instance, the matrix below has shape (3, 5)

```
> x <- matrix(rep(0, 3*5), nrow = 3, ncol = 5)
```

```
> x
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	0	0	0	0	0
[2,]	0	0	0	0	0
[3,]	0	0	0	0	0

```
> dim(x)
```

```
[1] 3 5
```

- The 3D tensor example has shape (2, 3, 2)

```
> x <- array(rep(0, 2*3*2), dim = c(2,3,2))
```

```
> str(x)
```

```
num [1:2, 1:3, 1:2] 0 0 0 0 0 0 0 0
```

```
> dim(x)
```

```
[1] 2 3 2
```

- A vector has a shape with a single element, such as (5).
- You can access the dimensions of any array using the `dim()` function.

# Data Type

- Type of the data contained in the tensor
  - For instance, a tensor's type could be integer or double
  - On rare occasions, you may see a character tensor



# Getting started with Neural Networks using Keras

# Workflow

1. Data Preprocessing
2. Constructing the model
3. Compile
4. Fit The Model (Train the model)
5. Evaluating your Model
6. Predict Labels of new data
7. Fine-tuning your model
8. Saving, Loading or Exporting your model

# 1) Data Preprocessing

- Reshaping and Scaling Data

**NOT PART OF THIS PRESENTATION**

## 2) Build the model

- 1) Define your **training** data: input tensors and target tensors.
- 2) Define a network of layers (or **model**) that maps your inputs to your targets.

`keras_model_sequential()`

```
model <- keras_model_sequential() %>%  
  layer_dense(units = 16,  
              activation = "relu",  
              input_shape = ncol(x_train)) %>%  
  layer_dense(units = 16,  
              activation = "relu") %>%  
  layer_dense(units = 1,  
              activation = "sigmoid")
```

### MODEL DEFINITION

What are my layers  
and how they are going  
to behave?

# Activation Functions

# Activation Function for each Layer

Example of activation functions: Sigmoid, Tanh and Relu

It appears that in most cases we can use the **relu** activation function. Relu looks like this:

- It is important to note that there's no **best** activation function. One may be better than other in many cases, but will be worse in some other cases.
- Another important note is that using different activations function doesn't affect what our network can learn, only how fast (how many data/epochs it needs).

# 3) Compile the model

What are the ***loss function*** and ***optimizer*** that I want to use during training and what **metrics** do I want to collect?

`compile()`

```
model %>% compile(  
  optimizer = "rmsprop",  
  loss = "binary_crossentropy",  
  metrics = c("accuracy")  
)
```

Note that the model is modified in place. It is all done in place and you don't assign to an object

## 4) Fit the model

Feeding mini-batches of data to the model thousands of times

`fit()`

```
model %>% fit(  
  x_train,  
  y_train,  
  batch_size = 512,  
  epochs = 4  
)
```

- Feed 512 samples at a time to the model (`batch_size = 512`)
- Transverse the input dataset 4 times (`epochs = 4`)



# 4) Fit the model


```
model %>% fit(  
  x_train,  
  y_train,  
  batch_size = 512,  
  epochs = 4,  
  validation_split = 0.2  
)
```

As we are fitting, we are going to hold out 20% of the data to **validate** that we are not just overfitting to our dataset

A deep learning model can memorize the data and it can give you a function that is not useful since what it just did was memorize the data

- Hold out 20% of the data to validation (`validation_split = 0.2`)

```
history <- model %>%  
  fit(x_train,  
    y_train,  
    batch_size = 512,  
    epochs = 4,  
    validation_split = 0.2)  
plot(history)
```



# 5) Evaluate the model

Feeding mini-batches of data to the model thousands of times

`evaluate()`

```
model %>% evaluate(x_test, y_test)
```

```
$loss
```

```
[1] 0.2916625
```

```
$acc
```

```
[1] 0.884
```

# 6) Predict Labels of new data

Generate predictions from the model

`predict()`

```
mmodel %>% predict(x_test[1:10,])
```

```
      [,1]  
[1,] 0.1922472  
[2,] 0.9998888  
[3,] 0.8777745  
[4,] 0.8634432  
[5,] 0.9594250  
[6,] 0.9000725  
[7,] 0.9998257  
[8,] 0.0138099  
[9,] 0.9733002  
[10,] 0.9950176
```